

Claude Code Workflow: Practical Patterns for the CLI Agent

How to get the most out of Claude Code as a solo dev — the memory system, plan mode, sub-agents, custom slash commands, hooks, and the prompting patterns that work. Written from inside this very setup, where Claude Code is the tool used to maintain the rest of these how-tos.

NOTE. LAST VALIDATED: 2026-05. Claude Code current (Opus 4.7 + Sonnet 4.6 + Haiku 4.5 model lineup), the memory + skill + hooks systems, MCP integration patterns. Bump and re-validate as the tool evolves — Claude Code is moving fast.

What you'll have at the end

- A clear mental model of what Claude Code is and how it interacts with your filesystem, shell, tools, and memory.
- A working `~/.claude/CLAUDE.md` (global) + per-project `CLAUDE.md` files that calibrate Claude's behavior without you having to re-explain context every session.
- A handful of **custom slash commands** for the workflows you repeat.
- **Hooks** for automated behaviors that should happen reliably (formatting, linting, notifications).
- A practical sense of **when to use plan mode**, **when to spawn sub-agents**, and **when to just chat**.
- Prompting patterns that produce useful output instead of generic mush.

Prerequisites

- Claude Code installed (`npm i -g @anthropic-ai/claude-code` or via your platform's installer).
 - An Anthropic API key or Claude.ai Pro/Max subscription with usage included.
 - Read access to `~/.claude/` and your project directories.
-

Table of contents

1. Mental model: what Claude Code actually is
2. The memory system: CLAUDE.md and auto-memory
3. Plan mode: when to slow down
4. Sub-agents: when to fan out
5. Slash commands and skills
6. Custom slash commands
7. Hooks: automated behaviors that don't require remembering
8. Permission management
9. Prompting patterns that work

10. [Working with dev containers](#)
 11. [MCP servers: extending capabilities](#)
 12. [Token budgets and context discipline](#)
 13. [Common workflows](#)
 14. [Alternatives considered](#)
 15. [Quick reference](#)
-

1. Mental model: what Claude Code actually is

Claude Code is a **command-line agent** that runs locally on your machine. It's not a chat client that calls an API in the abstract; it has direct access to:

- **Your filesystem** — can read, write, edit files in your working directory (and anywhere it has permission).
- **Your shell** — can run commands, see their output, react to results.
- **Built-in tools** — Read, Write, Edit, Bash, Grep, plus dozens more (TaskCreate, ExitPlanMode, SendUserFile, etc.).
- **Skills** — pre-built workflows you can invoke (`/review`, `/security-review`, `/init`, etc.).
- **MCP servers** — external tools that extend its reach (Gmail, Drive, Slack, custom).
- **Memory** — persistent context across sessions, both manual (`CLAUDE.md` files) and automatic (the `memory/` subsystem).

When you type a message, Claude can call any of these as part of formulating its response. The transcript you see is a conversation interleaved with tool calls, tool results, and Claude's text output.

[Diagram: see the online version on GitHub for the interactive flowchart.]

The agent loop runs until Claude decides it's done responding, then control returns to you for the next prompt.

What it isn't

- **Not a magic box** that writes whole projects from scratch reliably. It works in a tight loop with you; the more context you give about what good looks like, the more useful it is.
 - **Not stateless.** The current session has full context; across sessions, only what's in memory persists (`CLAUDE.md` and the memory subsystem).
 - **Not deterministic.** Same prompt twice may produce different outputs. Treat it as a smart colleague, not a build tool.
-

2. The memory system: CLAUDE.md and auto-memory

Memory is the single highest-leverage feature for solo work. Three tiers:

Tier 1: Global `CLAUDE.md` (`~/.claude/CLAUDE.md`)

Loaded into context for every Claude Code session, regardless of directory. Use for:

- Your environment (OS, shell, primary directories).
- Default tooling preferences (`uv` not `pip`, `bun` not `npm`, etc.).
- Personal background that influences calibration (CS background, time constraints, etc.).
- Non-negotiable preferences ("don't add Code Comments unless asked," "always use markdown for docs").

IMPORTANT. GLOBAL CLAUDE.MD IS LOADED INTO EVERY CONTEXT. Be selective — it costs tokens forever. ~150 lines is plenty. Anything topic-specific belongs in per-project `CLAUDE.md` files.

Tier 2: Project `CLAUDE.md` (in a repo)

Loaded when Claude Code is invoked inside that repo (or any subdirectory). Use for:

- Project-specific conventions (the writing standards in this very repo's `CLAUDE.md`, for instance).
- Architecture notes.
- "How to do X in this codebase" patterns.
- Things you'd otherwise re-explain every session.

There's a hierarchy: project `CLAUDE.md` adds to global `CLAUDE.md`. Both are in context.

Tier 3: Auto-memory (`~/.claude/projects/<...>/memory/`)

Claude writes memories automatically when you tell it personal facts, preferences, or project context. These are loaded into context too, indexed by `MEMORY.md`.

Categories:

- `user_*.md` — facts about you (role, background, expertise).
- `feedback_*.md` — guidance you've given (how you prefer work done, things to avoid/repeat).
- `project_*.md` — context about specific ongoing projects.
- `reference_*.md` — pointers to external resources.

The auto-memory system is bidirectional. Claude reads existing memory at session start; it writes new memory when something memorable comes up. You can also ask: "remember that I prefer X" → Claude writes a memory file.

TIP. AUDIT YOUR MEMORY OCCASIONALLY. `~/.claude/projects/<your-home-path>/memory/MEMORY.md` is the index. Open it, read the entries, delete what's stale. Memory drift is real — old preferences accumulate and start making advice worse.

What goes where (rule of thumb)

Content	Goes in
"I use Bun not Node"	Global CLAUDE.md
"This codebase uses SvelteKit + Drizzle"	Project CLAUDE.md
"I'm in the Army, doing an MSDS, time-constrained"	Auto-memory (user)
"Don't add try/except around code that can't fail"	Auto-memory (feedback)
"We're rewriting the auth system in Q3"	Auto-memory (project)
"Linear tickets live at company.linear.app/team/X"	Auto-memory (reference)

3. Plan mode: when to slow down

Plan mode is a special mode where Claude **doesn't make changes** — it only investigates and produces a written plan. You then approve (or revise) before any code is touched.

When to use plan mode

- **Multi-step changes** where you want to review the approach before committing time.
- **Anything destructive or hard to reverse** (data migrations, refactors that touch many files).
- **Cross-cutting concerns** where the impact isn't obvious upfront.
- **Estimation tasks** ("how big is this change?" — you want the plan as the answer, not the implementation).

When NOT to use plan mode

- Quick fixes (typo, single-file change, small refactor).
- Investigations / questions ("what does this file do?").
- Anything where you're going to immediately tell Claude to go anyway.

How to use it

```
/plan
```

Or just say "make a plan first" in your prompt. Claude switches mode, investigates, writes a plan, presents it via `ExitPlanMode`. You approve or reject; if you approve, Claude proceeds.

The plan should be specific enough that you could hand it to another engineer and they'd execute it. If a plan is vague, push back: "be more specific about what changes to which files."

4. Sub-agents: when to fan out

Claude Code supports **sub-agents** — spawning another Claude in a separate context to handle a discrete task, then returning the result to the main agent.

Available agent types:

- **general-purpose** — flexible; default if you don't specify.
- **Plan** — software architect; designs implementation strategies. Read-only.
- **Explore** — fast read-only search; for "find me X across this codebase."
- **claude** — same as the parent but in a fresh context. Useful for isolating noisy operations from the main context.
- **claude-code-guide** — Q&A about Claude Code itself.
- Other specialized agents (varies by version).

When to spawn a sub-agent

- **Big read-only research** that would clutter the main context. "Find every place we use `process.env.DATABASE_URL` and summarize." → spawn **Explore**.
- **Parallel independent work** — two unrelated investigations can run as two sub-agents simultaneously.
- **Heavy planning** where you want a fresh perspective. "Plan how to migrate from REST to GraphQL." → spawn **Plan**.
- **Anything that'd consume significant context but only the conclusion matters.**

When NOT to

- Single-file or single-step work. The overhead isn't worth it.
- Tasks where Claude needs to see your full context.
- Anything iterative or conversational.

Hot tip

Brief the agent like a smart colleague who walked in mid-meeting. They have no idea what the conversation has been. Give them:

- What you're trying to accomplish.
- What you've ruled out.
- The specific question or task.
- The format you want back ("under 200 words," "as a list," "as a code patch").

```
Agent({
  description: "audit cross-file consistency",
  subagent_type: "Explore",
  prompt: "We just renamed the `User` table to `Account` in src/db/schema.ts. Find every other file that references the old name (imports, queries, types, comments). Report a list of file:line locations. Under 300 words."
})
```

5. Slash commands and skills

Slash commands are pre-built workflows you can invoke. The most useful built-ins:

Command	What it does
<code>/help</code>	List available commands
<code>/init</code>	Generate a <code>CLAUDE.md</code> for the current project
<code>/review</code>	Review a PR (open or local branch)
<code>/security-review</code>	Security-focused review of the current branch
<code>/clear</code>	Clear the conversation context (start fresh)
<code>/compact</code>	Summarize old context to save tokens
<code>/agents</code>	List available sub-agents
<code>/permissions</code>	Manage tool permissions
<code>/memory</code>	View/edit memory entries

User-invocable skills (varies by setup) include things like:

- `update-config` — manage `settings.json`
- `keybindings-help` — keyboard shortcuts
- `simplify` — review code for quality
- `fewer-permission-prompts` — generate an allowlist from your transcripts
- `loop` — recurring task on an interval
- `schedule` — recurring cron-style remote agents

Type `/` to see what's available in your setup.

6. Custom slash commands

Define your own slash commands as markdown files in `~/.claude/commands/`. Each file is a command template.

Example: `~/.claude/commands/deploy-check.md`

```
You are checking whether the current branch is ready to deploy.

Run these checks in order:

1. `git status` – is the working tree clean?
2. `git diff origin/main..HEAD --stat` – what's changed vs main?
3. Run the test suite (`uv run pytest` or `bun test`, whichever is appropriate).
4. Run lints (`uv run ruff check` or `bun run lint`).
5. Check if `pyproject.toml` / `package.json` versions are bumped if applicable.

For each check, report: ☐ pass /  warn /  fail with one line of context.

End with an overall ship/don't-ship verdict.
```

Now `/deploy-check` runs that script-like prompt. Claude interprets each step as a goal and uses tools accordingly.

Other useful custom commands to write:

Command	Purpose
<code>/spike <topic></code>	Quick investigative spike — read related files, summarize the lay of the land
<code>/changelog</code>	Generate a changelog entry from recent commits
<code>/pr-description</code>	Write a PR description from the diff
<code>/migrate-todo <from> <to></code>	Walk through migrating from one tool/lib/version to another
<code>/dora</code>	Compute personal DORA metrics for your repos

Args

Pass arguments to commands by referencing `$ARGUMENTS` (or just embedding instructions for the user to provide them):

You are searching for occurrences of "\$ARGUMENTS" across the codebase.
Use Grep with sensible patterns; report file:line matches grouped by file.

Invoke: `/find-occurrences DATABASE_URL`.

7. Hooks: automated behaviors that don't require remembering

Hooks let `settings.json` define **shell commands that run on Claude Code events**. Examples:

- "After every file edit, run `prettier` on the edited file."
- "Before Claude reads a file, log the access for audit."
- "When the model stops talking, send a desktop notification."

Hooks are configured in `~/.claude/settings.json` or per-project `.claude/settings.json`. Schema (simplified):

```
{
  "hooks": {
    "PostToolUse": [
      {
        "matcher": "Edit|Write",
        "hooks": [
          {
            "type": "command",
            "command": "prettier --write '$CLAUDE_FILE_PATHS' 2>/dev/null || true"
          }
        ]
      }
    ],
    "Stop": [
      {
        "hooks": [
          {
            "type": "command",
            "command": "notify-send 'Claude finished'"
          }
        ]
      }
    ]
  }
}
```

When to use a hook

Anywhere you'd say "from now on, when X happens, do Y" — that's a hook, not a preference. Memory can store preferences but can't execute. Hooks execute.

Examples worth wiring up:

- **Auto-format after edits** — Prettier for TS, Ruff for Python, gofmt for Go.
- **Auto-lint warning** — fire ESLint/Ruff after edits, print warnings.
- **Notification on Stop** — desktop notification when Claude finishes responding (so you can wander off and come back).
- **Pre-flight check before Bash** — log destructive commands.

Skills for hooks

`/update-config` is the skill that helps you wire up hooks. Use it; don't hand-edit `settings.json` directly unless you're confident.

8. Permission management

Each tool Claude tries to use either:

- Auto-runs (whitelisted in `settings.json`'s `permissions.allow`).
- Prompts you (default for most tools).
- Auto-denies (blacklisted in `permissions.deny`).

The `fewer-permission-prompts` skill

Scans your recent transcripts for tools Claude has used repeatedly with your approval, then proposes adding them to the allowlist. Run it after a few weeks of use to dramatically cut the prompting.

```
/fewer-permission-prompts
```

It'll suggest entries like:

```
"permissions": {
  "allow": [
    "Bash(git status:*)",
    "Bash(git diff:*)",
    "Bash(ls:*)",
    "Bash(cat:*)",
    "Read",
    "Grep",
    "WebSearch"
  ]
}
```

Accept and you stop getting prompted for read-only operations.

What to never auto-allow

- **Anything that writes to remote services** (`git push`, `docker push`, `gh ...`) — interactive confirmation is the whole point.
- **Destructive operations** (`rm`, `dd`, `truncate`).
- **Anything that creates billable resources** (cloud APIs).

A reasonable allowlist is: read-only operations everywhere; writes prompt every time.

9. Prompting patterns that work

Be specific

Bad: "fix the auth bug"

Good: "the `/api/login` endpoint returns 200 with no body when password is correct but the session cookie isn't set. I think it's the cookie domain config. Trace the request through `src/lib/server/auth.ts` and `hooks.server.ts` and propose a fix."

Provide pointers, not synthesis

Bad: "based on your research, fix the bug"

Good: "the bug is in `src/lib/server/auth.ts:42`. The issue is that `secure: true` is set even in dev. Fix it conditionally based on `NODE_ENV`."

If you find yourself writing "based on the findings, do X" — you're abdicating the synthesis step. The synthesis step is where the value is. Do it yourself or have Claude make the plan first (plan mode).

State the constraints, including what NOT to do

"Implement the leaderboard endpoint. **Don't:** add new dependencies, change the existing schema, or modify the route file path."

Claude takes prohibitions seriously. Use them when needed.

Set output expectations

- "Reply in under 100 words."
- "Just write the code; don't explain it."
- "Output the diff only."
- "Summarize in 3 bullets."

Less constraint = more output. Pick the constraint that matches what you actually want.

Give worked examples for non-obvious style

If you have specific style preferences ("use guard clauses, not nested ifs" / "use `const` arrow functions, not `function`"), give a code snippet of the desired style. It's faster than describing.

When prompting for code: name the failure modes

"Write a function to parse the CSV. Handle these edge cases explicitly:

- Empty file → return `[]`.
- File with only a header row → return `[]`.
- Quoted strings with embedded commas → preserve.
- Missing trailing newline → still parse the last row.
- BOM at start → strip and parse normally.

Don't catch exceptions broadly; let parse errors propagate."

You get cleaner code because the spec is clearer.

10. Working with dev containers

Claude Code can run inside a Dev Container. The "Reopen in Container" flow attaches a remote Claude instance to the container environment. It uses the container's filesystem, the container's installed tools, and the SSH agent forwarded into the container (`ssh-keys`).

Why use Claude Code inside the container

- The tools Claude uses (uv, bun, alembic, pytest) are the project's tools — installed in the container, not on your host.
- Git operations Claude initiates use your forwarded SSH agent, so `git push` works without setup.
- Side effects (file writes, command output) stay scoped to the container.

Things to watch for

- `.claude/` in the project — settings and memory may or may not be in the repo. If you want them committed (project-specific), put them under `.claude/`. If not, leave them in `~/.claude/`.

- **Memory paths** are relative to `$HOME` of whatever user Claude is running as inside the container.
-

11. MCP servers: extending capabilities

MCP (Model Context Protocol) servers add tools beyond what Claude Code ships with. Common ones:

- **Gmail** — search threads, read messages, create drafts.
- **Drive** — search, read, create files.
- **GitHub** — issues, PRs (beyond what `gh` CLI exposes).
- **Slack** — read/write messages.
- **Linear / Notion / Asana** — task management integration.
- **Custom** — anything you can wrap in an MCP server.

Configure in `~/.claude/settings.json` under `mcpServers`. Each server runs as a subprocess; Claude calls its tools through the protocol.

TIP. MOST MCPS ARE LOADED LAZILY. They show up in the deferred-tools list and only get full schemas when Claude invokes `ToolSearch`. This keeps context-overhead low for unused integrations.

The Claude.ai connectors (Gmail, Drive, Calendar) are pre-wired if you've enabled them in your Claude.ai account settings.

12. Token budgets and context discipline

Each session has a context window (1M tokens for Opus 4.7 with 1M variant; 200k for others). You can lose effectiveness toward the end of long sessions when context fills up.

Strategies

- `/clear` — wipe the conversation, start fresh. Use after a big task is done.
- `/compact` — summarize the older portion of the conversation, retain recent specifics. Less brutal than `/clear`.
- **Spawn sub-agents for big read-only tasks** — keeps the main context lean.
- **Avoid asking Claude to "read all the files in the repo"** — it loads them into your context. Use Explore or Grep to surface only what you need.

Memory vs context

- **Context** = current conversation; gone after `/clear`.
- **Memory** = persisted across sessions in `~/.claude/`.

Move things from context → memory deliberately when they're stable preferences. The "remember that" prompt does this explicitly.

13. Common workflows

"Write a new feature end-to-end"

1. Open the repo in VS Code, reopen in dev container.
2. Type: "I want to add a `/leaderboard` slash command. Tables: there's already a `scores` table. UX: top 10 by total score. /plan first."
3. Claude enters plan mode, investigates, produces a plan.
4. You approve or revise: "look good — also add tests for empty leaderboard case."
5. Approve final plan.
6. Claude implements: writes the route, the test, the migration if needed.
7. You review the diff: `git diff`.
8. Push + PR via `gh pr create --fill`.

"Investigate a bug"

1. "There's a 500 error on `/api/auth/login` when password contains an emoji. Trace what happens."
2. Claude reads the relevant files, traces the call path, reports findings.
3. You decide on the fix (the value-add step) and tell Claude to implement it.

"Code review"

`/review` against the current branch. Claude reads the diff, comments on what it sees, suggests improvements. Iterate or accept.

"Refactor"

For multi-file refactors, plan mode is essential. Without it, Claude can drift mid-task. With it, you've agreed on the approach before edits start.

"Write or update a how-to"

Following this repo's pattern: write the README, update CLAUDE.md and the top-level index, commit + push. Claude can drive this loop with one or two prompts if the topic is clear.

14. Alternatives considered

Tool	Verdict	Why
Claude Code (this guide)	Default.	Best agent UX as of 2026; memory + skills + hooks are unique; great Anthropic model lineup.
Cursor	Solid editor-integrated alternative.	If you want IDE-first rather than CLI-first. Editor-native UX wins for some workflows.
Aider	Open-source CLI agent.	More transparent (you see every call), works with multiple LLM backends. Less polish than Claude Code.
GitHub Copilot Chat	Built into VS Code.	Better for "this line, suggest a completion." Worse for "make a multi-step change across the codebase."
Continue	Open-source VS Code extension.	Multi-LLM, configurable. Smaller community than Cursor.
Plain ChatGPT/ Claude.ai web	Fine for one-off questions.	No filesystem access, no shell, no memory persistence with your real code. Lose the agent advantages.
Sourcegraph Cody	Enterprise-focused.	Better for organization-wide code search; overkill for solo.

For solo dev workflow, **Claude Code wins on the trifecta of memory + agentic loop + Anthropic model quality**. The cost is being tied to one vendor; the gain is the fastest path from "I have an idea" to "PR is open."

15. Quick reference

Files and dirs

Path	Purpose
<code>~/.claude/CLAUDE.md</code>	Global instructions, loaded every session
<code>~/.claude/settings.json</code>	Permissions, hooks, MCP servers
<code>~/.claude/commands/*.md</code>	Custom slash commands
<code>~/.claude/projects/<...>/memory/MEMORY.md</code>	Auto-memory index
<code>~/.claude/projects/<...>/memory/*.md</code>	Individual memory files
<code><repo>/CLAUDE.md</code>	Project-specific instructions
<code><repo>/.claude/settings.json</code>	Project-specific settings

Slash commands worth knowing

```
/init          # generate CLAUDE.md for current repo
/clear         # wipe conversation context
/compact      # summarize old context
/plan         # enter plan mode
/review       # PR review of current branch
/agents       # list available sub-agents
/permissions  # manage tool allowlist
/memory       # view/edit memory
/help         # list everything
/fewer-permission-prompts # auto-generate allowlist from transcripts
/update-config # configure hooks/settings
```

Custom command template

You are <doing a thing>.

Steps:

1. ...
2. ...

Constraints:

- Don't ...
- Do ...

Output format:

- ...

Save as `~/.claude/commands/<name>.md`, invoke as `/<name>`.

Hook templates

`settings.json`:

```
{
  "hooks": {
    "PostToolUse": [
      { "matcher": "Edit|Write", "hooks": [
        { "type": "command", "command": "prettier --write '$CLAUDE_FILE_PATHS' 2>/dev/
null || true" }
      ]
    },
    "Stop": [
      { "hooks": [
        { "type": "command", "command": "notify-send 'Claude finished'" }
      ]
    }
  ]
}
```

Prompting cheat sheet

- **Be specific.** Names, file paths, function names, expected behavior.
- **State constraints.** What NOT to do is just as important as what TO do.

- **Set output format.** "Reply in 100 words." "Diff only." "Bullet list."
- **Don't abdicate synthesis.** If you find yourself saying "based on findings, implement X," do the synthesis yourself.
- **Use plan mode** for multi-step or destructive work.
- **Use sub-agents** for big read-only exploration or parallel independent work.

When to /clear

- After completing a big task and starting an unrelated one.
- When responses start getting slower or less relevant.
- When the conversation has accumulated 50+ messages of mixed topics.

Memory hygiene

- `/memory` to view what's been captured.
- Read `MEMORY.md` index quarterly; delete stale entries.
- Be explicit when you want something remembered: "remember that I prefer X."
- Be explicit when something is one-off: "for this conversation only..."